

1 Week1

Introduction to Deep Learning.

1.1 What is a Neural Network?

1.2 Supervised Learning with Neural Networks

1.3 Why is Deep Learning taking off?

1.4 About this Course


1. Neural Networks and Deep Learning
2. Improving Deep Neural Networks: Hyperparameter tuning, Regularization and Optimization
3. Structuring your Machine Learning project
4. Convolutional Neural Networks
5. Natural Language Processing: Building sequence models

1.5 Outline of this Course

- Week1. Introduction
- Week2. Basics of Neural Network programming
- Week3. One hidden layer Neural Networks
- Week4. Deep Neural Networks

2 Week2

Basics of Neural Network Programming
How do I write an equation in L^AT_EX?



In 1902, Einstein created this equation: $E = mc^2$
And Newton came up with this one: $\sum F = ma$

$$5 + 5 = 10 \tag{1}$$

$$\begin{aligned} A &= \frac{5\pi r^2}{2} \\ A &= \frac{1}{2}\pi r^2 \end{aligned} \tag{2}$$

2.1 Neural Network Notations

General comments:

superscript (i) will denote the i^{th} training example.

superscript $[l]$ will denote the l^{th} layer.

Sizes:

- m : number of examples in the dataset
- n_x : input size
- n_y : output size
- $n_h^{[l]}$: number of hidden units of the l^{th} layer.
In a for loop, it is possible to denote $n_x = n_h^{[0]}$ and $n_y = n_h^{[numberoflayer+1]}$
- L : number of layers in the network

Objects:

- $X \in \mathbb{R}^{n_x \times m}$ is the input matrix
- $x^{(i)} \in \mathbb{R}^{n_x}$ is the i^{th} example represented as a column vector
- $Y \in \mathbb{R}^{n_y \times m}$ is the label matrix
- $y^{(i)} \in \mathbb{R}^{n_y}$ is the output label for the i^{th} example
- $W^{[l]} \in \mathbb{R}^{numberofunitsinnextlayer \times numberofunitsinthepreviouslayer}$ is the weight matrix, superscript $[l]$ indicates the layer
-

Common forward propagation equation examples:

-
-

Examples of cost functions:

-
-

2.2 Binary Classification

Use matrix without using for loops.

Computation using Forward propagation and Backward propagation.

Logistic regression is an algorithm for binary classification.

m training examples $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
where $x^{(i)} \in \mathbb{R}^{n_x}$ and $y^{(i)} \in \{0, 1\}$ for $i \in [1, m]$

$$X = \begin{bmatrix} \vdots & \vdots & \vdots \\ X^{(1)} & X^{(1)} & X^{(m)} \\ \vdots & \vdots & \vdots \end{bmatrix} \in \mathbb{R}^{n_x \times m}$$

$$X.shape = (n_x, m)$$

$$Y = [Y^{(1)}, Y^{(2)}, \dots, Y^{(m)}] \in \mathbb{R}^{1 \times m}$$

$$Y.shape = (1, m)$$

2.3 Logistic Regression


Given X , want $\hat{Y} = P(Y = 1|X)$ where $X \in \mathbb{R}^{n_x}$

Parameters: $\omega \in \mathbb{R}^{n_x}$ a n_x dimensional vector, $b \in \mathbb{R}$ a real number.

Output $\hat{y} = \sigma(\omega^T X + b) = \sigma(z)$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Drawing a sigmoid function and its derivative in tikz



2.4 Logistic Regression Cost Function

To train the parameter ω and b of a Logistic Regression Model, we need a cost function.

$$\hat{y} = \sigma(\omega^T X + b) \text{ where } \sigma(z) = \frac{1}{1+e^{-z}}$$

$$\hat{y}^{(i)} = \sigma(\omega^T X^{(i)} + b) \text{ where } \sigma(z^{(i)}) = \frac{1}{1+e^{-z^{(i)}}}$$

Given $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$, want $\hat{y}^{(i)} \approx y^{(i)}$.

Loss(error) function (for a single training Example):

$$\mathcal{L}(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log (1 - \hat{y}))$$

Cost function (for the entire training Examples):

$$\mathcal{J}(\omega, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log (1 - \hat{y}^{(i)})]$$

The loss function computes the error for a single training example; the cost function is the average of the loss functions of the entire training set.


In training logistic regression model, we will try to find ω and b such that they minimize the Cost function $\mathcal{J}(\omega, b)$.

Logistic Regression can be seen as a very small Neural Network.

2.5 Gradient Descent

$$\mathcal{J}(\omega, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log (1 - \hat{y}^{(i)})]$$

Want to find ω and b that minimize the Cost function $\mathcal{J}(\omega, b)$.



$\mathcal{J}(\omega, b)$ is a convex function with a single local optimum.

No matter where you initialize the point, you should get to the same point (Global optimum).

Repeat:

$$\omega := \omega - \alpha \frac{\partial \mathcal{J}(\omega, b)}{\partial \omega}$$

$$\omega := \omega - \alpha d\omega$$

$$b := b - \alpha \frac{\partial \mathcal{J}(\omega, b)}{\partial b}$$

$$b := b - \alpha db$$

where α is the learning rate.

2.6 Derivatives

derivatives; slope

Given $f(a) = 3a$


$\epsilon = .001, a = 2 + \epsilon$

$$\frac{f(a) - f(a + \epsilon)}{\epsilon}$$

make ϵ close to zero \rightarrow derivatives.

$$\frac{df(a)}{da} = \frac{d}{da} f(a)$$

2.7 More Derivative Examples



$$a = 2, f(a) = 4$$

$$a = 2.001, f(a) = 4.004001$$

$$\frac{d}{da} f(a) = 4, \text{ when } a = 2$$

$$\frac{d}{da} f(a) = 10, \text{ when } a = 5$$

$$\frac{d}{da} f(a) = \frac{d}{da} a^2 = 2a$$

Given a nudge $\epsilon = 0.001$ to a , the $f(a)$ goes up $2 * a$.

2.8 Computation Graph

- Forward propagation step(forward pulse) : compute output of the network.
- Backward pulse: compute the gradients or derivatives.

$$J(a, b, c) = 3(a + bc)$$

$$u = bc$$

$$v = a + u$$

$$J = 3v$$

In order to compute derivatives, you go backward propagation.

One step of backward propagation on a computation graph yields derivative of final output variable.

2.9 Computing derivatives.

$$u = bc$$

$$v = a + u$$

$$J = 3v$$

Given $a = 5, b = 3, c = 2$, then $v = 11, J = 33$

We want to see how much J changes if we change the values of a, b, c, u, v
for $J = 3v$

If we increase v to 11.001, then $J = 33.003$.

$$\frac{dJ}{dv} = 3$$

$$a = 5 \rightarrow a = 5.001$$

$$v = 11 \rightarrow v = 11.001$$

$$J = 33 \rightarrow J = 33.003$$

By Chain Rule:

$$\frac{dJ}{da} = 3 = \frac{dJ}{dv} \frac{dv}{da} = 3 \times 1$$

$$\frac{dFinalOutputVar}{dvar} \text{ where } var \text{ can be } a, b, c, \dots$$

We can simply denote $\frac{dJ}{dv} = dv$, $\frac{dJ}{da} = da$, etc with respect to J .

Similarly, $\frac{dJ}{du} = \frac{dJ}{dv} \frac{dv}{du} = 3 \times 1$

$\frac{dJ}{db} = \frac{dJ}{du} \frac{du}{db} = 3 \times 2 = 6$, where $u = bc = 2b$, and $\frac{du}{db} = 2$, given $a = 5, b = 3, c = 2$.

$\frac{dJ}{dc} = \frac{dJ}{du} \frac{du}{dc} = 3 \times 3 = 9$, where $u = bc = 3c$, and $\frac{du}{dc} = 3$, given $a = 5, b = 3, c = 2$.

$$\frac{dJ}{da} = 3$$

$$\frac{dJ}{du} = 3$$

$$\frac{dJ}{db} = 6$$

$$\frac{dJ}{dc} = 9$$

The coding convention $dvar$ represents: The derivative of a final output variable with respect to various intermediate quantities.

2.10